

setup for ollama service

ollama service kører på 11435 for ikke at kollidere med docker på 11434

```
# se loadede modeller
curl -s http://127.0.0.1:11435/api/tags | grep -o '"name": "[^"]*"'
"name": "gemma4:31b"
"name": "gemma4:26b"
"name": "gemma4:12b"
"name": "gemma4:latest"
```

gemma4 modellerne er sat op til at bruge 32k context

```
# model file
cat Modelfile-gemma4-31b-32k

FROM gemma4:31b
PARAMETER num_ctx 32768

# create model med 32k context
ollama create gemma4-31b:32k -f ~/src/ollama-service/Modelfile-gemma4-31b-32k

# run model
ollama run gemma4-31b:32k

# curl test
curl -X POST http://127.0.0.1:11435/api/generate -H "Content-Type: application/json" -d '{
  "model": "gemma4",
  "prompt": "Hej, svar med ordet test",
  "stream": false
}'

{"model": "gemma4", "created_at": "2026-06-16T06:20:16.909043332Z", "response": "Test", "done": true, ...

# check modeller loaded i ollama
ollama ps
```

NAME	ID	SIZE	PROCESSOR	CONTEXT	UNTIL
gemma4-31b:32k	36f5ca871cd0	29 GB	100% GPU	32768	24 hours
from now					
gemma4-26b:32k	8ab8e8b44290	19 GB	100% GPU	32768	23 hours
from now					
gemma4-latest:32k	89b70054091c	4.0 GB	100% GPU	32768	23 hours
from now					